Underline = wording?

Introduction

Nonprofits have always had to fight to make their voice heard; unless attached to a recognizable name, backed by corporate philanthropy, nonprofit organizations much of the time are playing from behind. The key differences between a nonprofit and a for-profit organization–though seemingly obvious from name alone–have become much more evident with the rise of artificial intelligence. AI as the new technology of today has provided companies with high rates of efficiency, but people are getting hired less and less. Compare this to a nonprofit whose "bottom line" is determined by the passion and care of the people behind it, whose core value is not defined by a net revenue but their contribution to a net good–one cannot have a nonprofit without the people. And as people become more of a rarity in our interactions with technology–chatbots dialing your phone number, texting and emailing you–we must consider how to confront the technology and make it work for *us,* not *for* us.

For-Profits and Nonprofits

This is the section where I would talk about grant writing and such but I haven't figured out how to stitch the pieces together yet

The AI Stigma

AI is of course the subject of countless debates in a variety of fields, with many palpable, valid concerns that should not be ignored. But though it may sound defeatist, there is a point at which principle may read as naivete. For-profits are not too interested in the high road, but the introduction of AI into workflow is no longer a possibility, it is a reality; agents powered by LLMs have reached a computational power previously unheard of, and to not jump at the opportunity would put one at a disadvantage. Nonprofit organizations, on the other hand, are relatively interested in the high road, with like-minded individuals rallying together to support a social effort–and Seeds of Success is no different. Therefore, there existed a need to shape the technology, a way to ensure that the organization doesn't fall behind while still maintaining credibility.

For many, proponents and skeptics alike, a major concern with the usage of LLMs is hallucination, especially with the widespread implementation of AI in search engines like Google; something presenting itself as a credible, knowledgeable source while providing incorrect, inaccurate information creates a loss of trust. This stems from a lack of information in

its knowledge base about a particular topic but the need to still produce an authoritative output, particularly if calibrated with a reward-based system that prioritizes output over correctness. So how do we maintain the efficiency of LLM-based agents and trust simultaneously? Seeds of Success's solution is the **living knowledge base**, a source of truth that represents the newest information composed of human-verified documents only related to the organization itself. This practice provides the agent with a direct highway to the information, decreasing resource use, response time, and increasing real output per dollar while maintaining that degree of trust; if AI is trained to pull only from what it knows it knows, contained within documents and files whose certainty have been determined by a human, hallucination is mitigated nearly entirely as assumptions and attempts at them are reduced.

And this couples directly with Seeds of Success's philosophy: use AI to help people, not replace them–"for access, not creation." Job security became another concern as AI entered the public eye; if the introduction of machinery and automation left many uncertain of the job market before, AI and its rapid growth in capability skyrocketed this to an unprecedented level as more and more jobs became perceived as expendable. But the integration of AI in nonprofit work does not seek to take the heart out of nonprofits, it seeks to empower the people within them. Through use of the knowledge base, less of a volunteer's time is spent teaching and more time is spent doing, enabling them to carry themselves forward through their own motivation without being bogged down by minutia. This was especially true for an organization like KARA, which after decades of collecting and storing data, statistics, articles, podcasts–it was hard not to get overwhelmed by sheer volume. **ELLA** sought to create an intuitive interface and surface the most pertinent information for any given task, facilitating work to be done by humans with the assistance of AI to retrieve the correct documentation. And as they produce, revise, iterate and add to the knowledge base, all changes are made by them and verified by them.

The Bottom Line is People

Much of the issues that stemmed from the rapid implementation of AI in our daily lives came about because companies forgot who they were building for; a product was built, deployed untested, improved exploitatively, marketed insincerely, and then used to push out the very group that provided their training data.

Take the recent controversy surrounding OpenAI's supposed proof of a longstanding Millennium Prize problem regarding the Navier-Stokes Equation: OpenAI asserted that a "system of coordinating agents" powered by an "internal model" was able to prove the breakdown–the existence of no solutions–of the equation given some initial conditions. And

though the specifics are outside of the scope of this paper, it rapidly caused controversy within the math community, especially for two mathematicians, Buckmaster and Alpöge, whose exact ideas seemed to be used with OpenAI's internal model to snatch the rug from under them. Furthermore, the proof–regardless of correctness–troubled many dedicated mathematicians, as it represented a treatment of these problems as a "benchmark," when "solving problems is only a tool" to achieve "conceptual understanding and insight." In other words, the black-box nature of the system itself stripped away the real prestige that came with solving the problem: the knowledge and foundation that naturally developed for others in doing the work.

This demonstrates a key insight that some fail to consider, which is that maybe people *do* want to do the work, not only for their own edification but for those that come after them as well. Nonprofits in particular can sympathize heavily with this mindset; volunteers for nonprofits are not often there to stay, but are still wanting to maximize the impact that they have on the particular sector they're working in and the organization itself. Attempting to outsource this passion to an LLM would be a disservice both to the individuals that possess it and the cause being supported. Seeds of Success aims to instead foster that passion, to give individuals the skills needed to compete in the job market while maintaining their work ethic; in 2026, 70% of employers use skill-based hiring practices, meaning education and credentials are no longer enough to reliably find a job, let alone land an interview. It also means that the skill ceiling has been raised a significant amount in recent years, especially given the rapid improvement of LLM-based agents. So in a volunteer's time with Seeds of Success, the hope is to **upskill** them, to combine their passion for their work with the ability to do all that it calls for. AI allows for the production of fast prototypes to demonstrate the viability of concepts, while humans are responsible for verification and actual product implementation. As a result, any gaps in a volunteer's knowledge can be filled for the time being, and–as long as they do their due diligence–they can become competent in creating products that people actually want to use to contribute to a positive social good.

These products can then be made **open source**, to let other volunteers in the future understand how they work and to iterate, improve upon them. This is the first move towards an open source future for nonprofits–with the beginning within Seeds of Success itself–to reduce nonprofit dependency on for-profit software when they have very few resources to spare to begin with.

Sources

(Not properly cited, just links for now)

https://web.archive.org/web/20260908171518/https://openai.com/index/navier-stokes-solution/

https://www.ibm.com/think/topics/ai-hallucinations

https://cims.nyu.edu/~tristanb/statement.pdf

https://zenodo.org/records/22737751

https://www.nu.edu/blog/67-hiring-statistics/